

PRINCESS NOURAH UNIVERSITY

DEPARTMENT OF COMPUTER AND INFORMATION SCIENCES

Masters of Data Science

First Term 2023



Breast Cancer

Data Management and Visualization

This report is submitted in full fulfillment of the Data Management and Visualization Project

Academic year: 2023/2024

Author:

Sawsan Daban 445009481

Alaa Alsharekh 445009444

Supervisor:

Dr. Romana Aziz

ACKNOWLEDGEMENTS

We express our deep gratitude to the Almighty Allah for giving us an opportunity to successfully finish modeling and simulating and write this report.

We would like to express our gratitude and appreciation to Princess Nourah Bint Abdulrahman University for giving us the opportunity to continue learning and achieve new things in our life, and for their remarkable services they gave us.

We would like to express our gratitude to Dr. Romana Aziz for her kindness and unequivocal support.

We would like to thank our friends for helping us realize what We are missing from our life. We thank them for their interesting discussions and help. Without them, We would not be here.

Finally, We would like to thank our family for their unwavering support throughout our lives, pushing us to always do our best.

EXECUTIVE SUMMARY

This report includes information about the problem statement of the project, goals and objectives, data preparation, and data visualization.

Contents

1	Introduction	1
1.1	Background	1
1.2	Problem Statement	1
1.3	Goals And Objectives	1
2	Data Preparation	2
2.1	Data Summary	2
2.1.1	Quality Report	3
2.2	Data Quality Control	5
2.2.1	Data consistency	6
2.2.2	Error Checking	9
2.2.3	Zero versus null	11
2.2.4	Validation step	12
3	Univariate Bivariate Analysis	13
3.1	Univariate Analysis	13
3.1.1	Qualitative variable	13
3.1.2	Quantitative variable	15
3.2	Bivariate Analysis	17
3.2.1	Categorical vs Quantitative	17
3.2.2	Quantitative vs Quantitative	19
	Appendices	20
A	Full Report	21

List of Figures

2.1	Quality report for row data	3
2.2	Quality report showing current data set issues	3
2.3	Quality report for Menopause feature	4
2.4	Modify data types	6
2.5	Check for duplication	7
2.6	Remove duplication	7
2.7	Check the data shape after removing duplication	7
2.8	Assign a name to the unnamed column	8
2.9	Modify column names to start with Capital letters	8
2.10	Dataset after removing duplicates	8
2.11	Error checking - typos	9
2.12	Error checking - typos	10
2.13	Handle null values	11
2.14	Validation step	12
2.15	Resolving dataset after validation	12
3.1	Univariate Analysis - Menopause	13
3.2	Univariate Analysis - Breast	14
3.3	Univariate Analysis - Age	15
3.4	Univariate Analysis - Nodes	16
3.5	Bivariate Analysis - Menopause vs Age	17
3.6	Bivariate Analysis - Menopause vs Tumor-size-encoded	18
3.7	Bivariate Analysis - Age vs Node	19
3.8	Bivariate Analysis - Age vs Tumor-size _{encoded}	19

List of Tables

None

Introduction

1.1 Background

Our Breast Cancer dataset appears to represent a sample of individuals, likely patients, from a medical context. These individuals are described based on various characteristics related to medical conditions.

Each row in the dataset represents an observation of an individual patient. There are multiple observations, each corresponding to a different patient in the dataset.

1.2 Problem Statement

The dataset comprises information about individuals diagnosed with breast cancer. It includes details such as age, menopausal status, tumor size, number of nodes, positivity or negativity of node-caps, degree of malignancy, affected breast, affected quadrant within the breast, receipt of irradiation treatment, and recurrence status of the cancer.

1.3 Goals And Objectives

- Acquire a comprehensive understanding of data preparation for analysis.
- Interpret data using graphical representations.
- Perform univariate and bivariate analysis using suitable data visualization techniques.

Data Preparation

2.1 Data Summary

The dataset comprises several numerical features that provide valuable insights into the characteristics of individuals in the study. Firstly, the "Age" variable represents the age of the subjects, offering a continuous numerical measure. Moving on to the "Tumor-size" feature, despite its apparent string format denoting intervals ('15-19', '35-39', etc.), it fundamentally represents numerical data, allowing for a quantitative assessment of tumor sizes. Lastly, the "Nodes" variable contributes discrete numerical values, indicating the count of nodes associated with each individual. Together, these numerical attributes form a foundation for quantitative analysis and interpretation in the dataset.

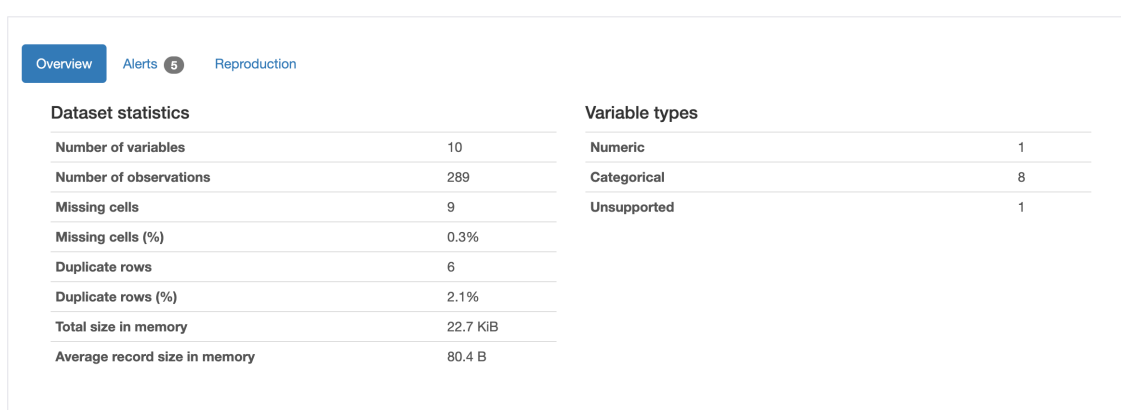
The dataset also incorporates several categorical features that play a pivotal role in characterizing individuals within the study. Firstly, the "Menopause" variable classifies individuals based on their menopausal status, with categories including 'premeno' and 'ge40'. Moving on to the "Degree-of-malignance," this categorical feature reflects the degree of malignancy, with categories such as '3', '1', and '2'. The "Breast" variable categorizes individuals based on the side of the breast affected, distinguishing between 'right' and 'left'. Similarly, the "Breast-quad" feature categorizes the quadrant of the breast affected, with categories like 'left_up' and 'central'. The "Irradiation" variable is binary, indicating whether irradiation was administered, with categories 'yes' or 'no'. Finally, the "Recurrence" variable serves as a categorical indicator of events, with potential values being 'recurrence-events' or 'no-recurrence-events'. Together, these categorical attributes offer a comprehensive understanding of qualitative aspects in the dataset.

2.1.1 Quality Report

- Overview ¹

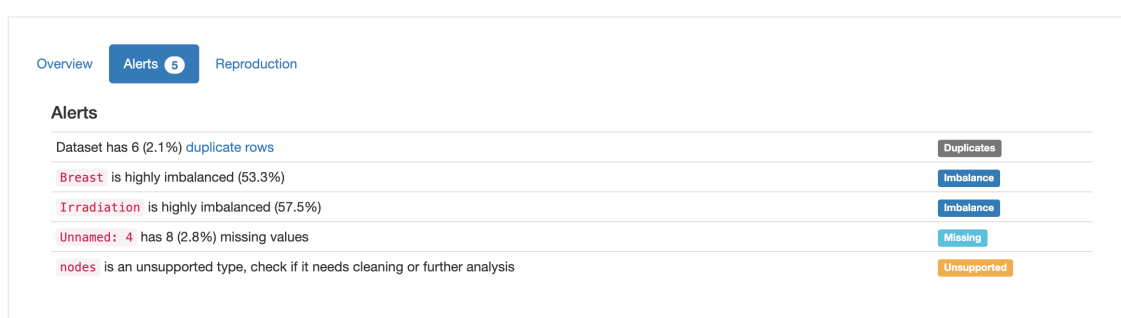
The provided summary offers a concise overview of the existing dataset and highlights certain noteworthy aspects. According to the description, our dataset comprises 10 variables and encompasses a total of 289 observations. Among these variables, 8 fall under the categorical category, while 1 is designated as a numerical variable in addition to unsupported variable due to unknown type .2.1 Additionally, the summary points out the presence of duplicate rows and an unnamed column as part of the dataset's characteristics.

2.2



Dataset statistics		Variable types	
Number of variables	10	Numeric	1
Number of observations	289	Categorical	8
Missing cells	9	Unsupported	1
Missing cells (%)	0.3%		
Duplicate rows	6		
Duplicate rows (%)	2.1%		
Total size in memory	22.7 KiB		
Average record size in memory	80.4 B		

Figure 2.1: Quality report for row data



Alerts		
Dataset has 6 (2.1%) duplicate rows		Duplicates
Breast is highly imbalanced (53.3%)		Imbalance
Irradiation is highly imbalanced (57.5%)		Imbalance
Unnamed: 4 has 8 (2.8%) missing values		Missing
nodes is an unsupported type, check if it needs cleaning or further analysis		Unsupported

Figure 2.2: Quality report showing current data set issues

¹For more information, please check YData Profiling

- Variables

As an additional feature in the quality report, a comprehensive description of each variable has been provided in the report as shown in the example below. 2.3

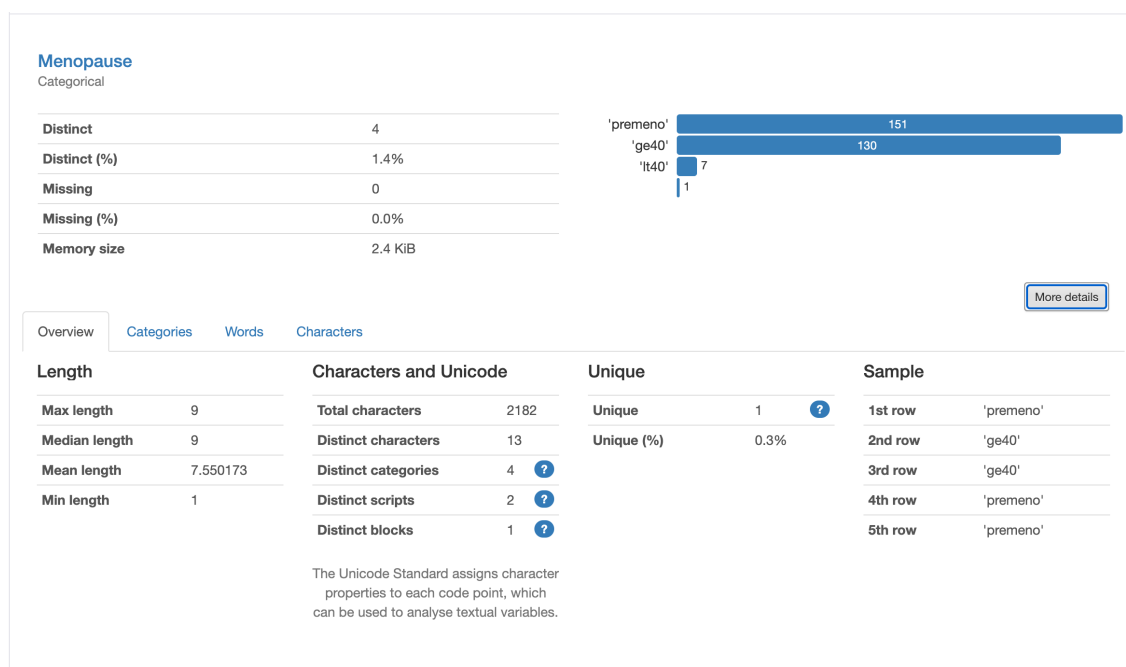


Figure 2.3: Quality report for Menopause feature

After conducting a quality report, the identified challenges and issues within the dataset have been addressed. This process involves implementing various data preprocessing techniques to enhance the quality and reliability of the data. By handling missing values, duplicates, outliers, and other inconsistencies, analysts can create a more accurate dataset for analyses. This, in turn, enables informed decision-making and contributes to the overall reliability of the results obtained from the data.

2.2 Data Quality Control

Every set of data needs to undergo a form of quality control, which includes verifying for errors and ensuring uniform consistency, this might include :

1. Data consistency
2. Error checking
3. Zero versus null

2.2.1 Data consistency

Ensuring data consistency involves modifying data types, examining the dataset for duplicate entries, confirming the presence of column names, and validating the conventions used for column naming.

- *Data type*

Converting the data type is an efficient pre-processing to handle the data in a better and more optimized way. This process ensures that each variable is assigned the most suitable data type.

During this step, the focus was on addressing unsupported variable types highlighted in the previous data quality report 2.2, with the aim of converting them into numeric types.

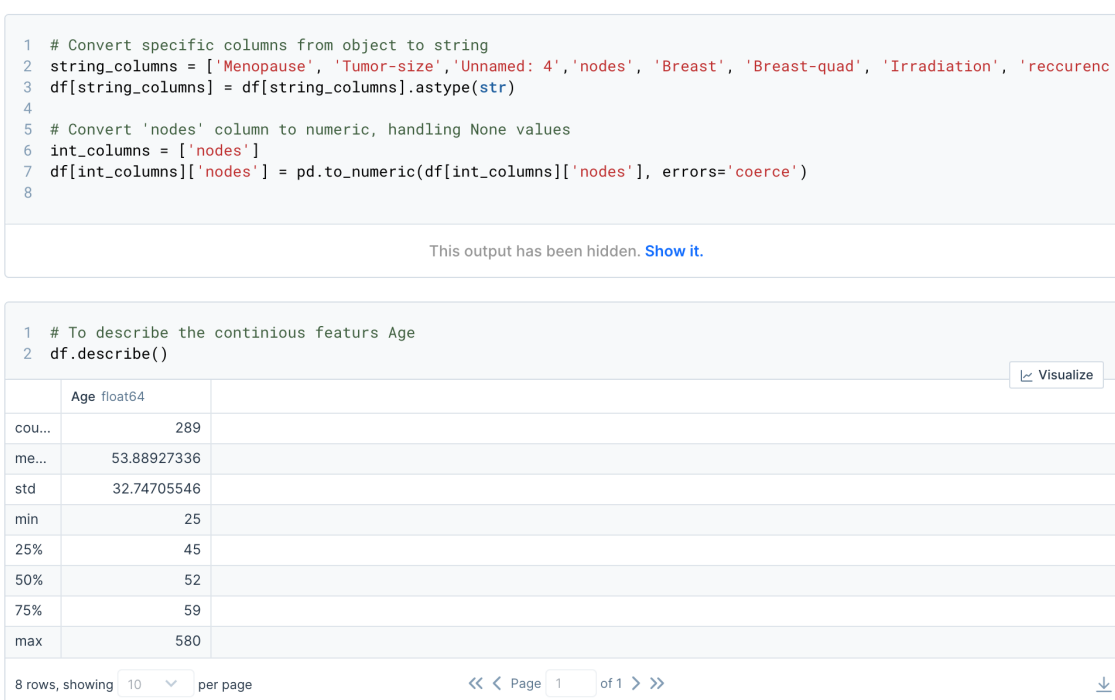


Figure 2.4: Modify data types

- *Duplication*

Initially, the dataset contained four duplicate rows. These were addressed by removing the duplicates.

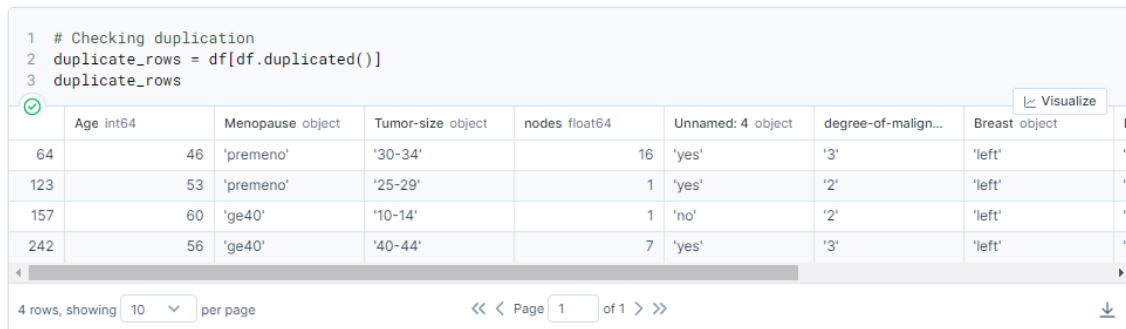


Figure 2.5: Check for duplication

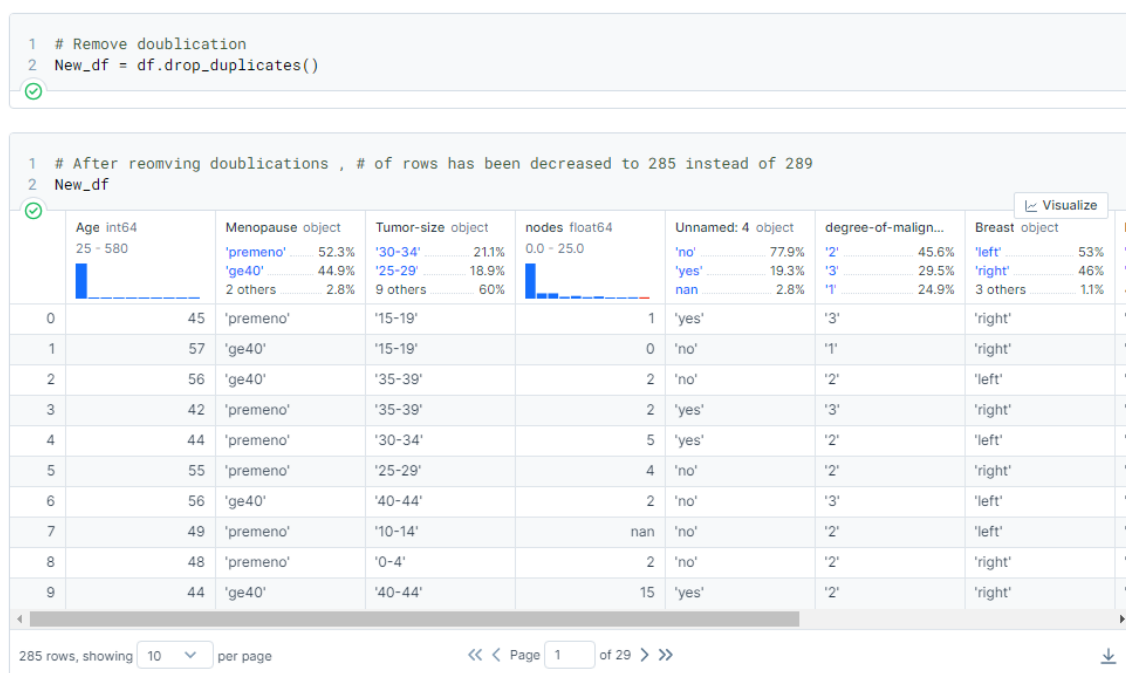


Figure 2.6: Remove duplication




Figure 2.7: Check the data shape after removing duplication

- *Missing Column Names*

It was noticed that one column was unnamed. In light of understanding the problem statement and the values associated with this feature, the decision was made to rename it to Node-caps.

```
1 # Assign a name to the unnamed column (Unnamed: 4)
2
3 New_df.rename(columns={'Unnamed: 4': 'Node-caps'}, inplace=True)
```

Figure 2.8: Assign a name to the unnamed column

An additional step has been added, by unifying the column to start with capital letters.

```
1 # Modify column names to start with Capital letters
2
3 New_df.rename(columns={'nodes': 'Nodes'}, inplace=True)
4 New_df.rename(columns={'degree-of-malignance': 'Degree-of-malignance'}, inplace=True)
5 New_df.rename(columns={'reccurrence': 'Reccurrence'}, inplace=True)
6
```

1 # After removing douplication & rename the columns
2 New_df

	Age int64 25 - 580	Menopause object 'premeno' 52.3% 'ge40' 44.9% 2 others 2.8%	Tumor-size object '30-34' 21.1% '25-29' 18.9% 9 others 60%	Nodes object 2 33% 1 31.2% 17 others 35.8%	Node-caps object 'no' 77.9% 'yes' 19.3% nan 2.8%	Degree-of-malign... '2' 45.6% '3' 29.5% '1' 24.9%	Breast object 'left' 53% 'right' 46% 3 others 1.1%
0	45	'premeno'	'15-19'	1	'yes'	'3'	'right'
1	57	'ge40'	'15-19'	0	'no'	'1'	'right'
2	56	'ge40'	'35-39'	2	'no'	'2'	'left'
3	42	'premeno'	'35-39'	2	'yes'	'3'	'right'
4	44	'premeno'	'30-34'	5	'yes'	'2'	'left'
5	55	'premeno'	'25-29'	4	'no'	'2'	'right'
6	56	'ge40'	'40-44'	2	'no'	'3'	'left'
7	49	'premeno'	'10-14'		'no'	'2'	'left'
8	48	'premeno'	'0-4'	2	'no'	'2'	'right'
9	44	'ge40'	'40-44'	15	'yes'	'2'	'right'

285 rows, showing 10 per page

Figure 2.9: Modify column names to start with Capital letters

```
1 New_df.shape
```

(285, 10)

Figure 2.10: Dataset after removing duplicates

2.2.2 Error Checking

In error checking, a verification process is implemented to examine feature values and make necessary adjustments.

- *Typos*²

As observed, certain features exhibit inconsistent values, such as Breast, Recurrence, and Irradiation. To address this issue, a unification of values is performed.

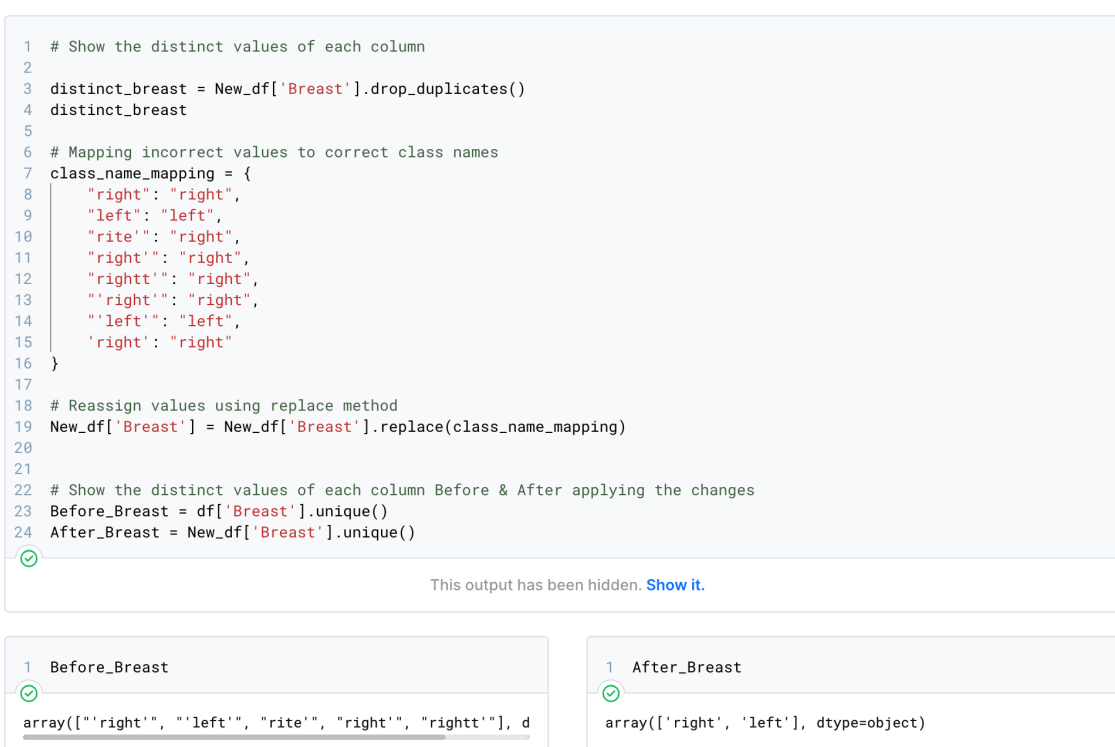


Figure 2.11: Error checking - typos

²The remaining features are added in the Appendices

- Outliers

The average Age is approximately 52. Excluding the outlier with the value 580, the maximum Age is 77, and the minimum Age is 25. The data anomaly associated with the value 580 has been identified and corrected as shown below. 2.12

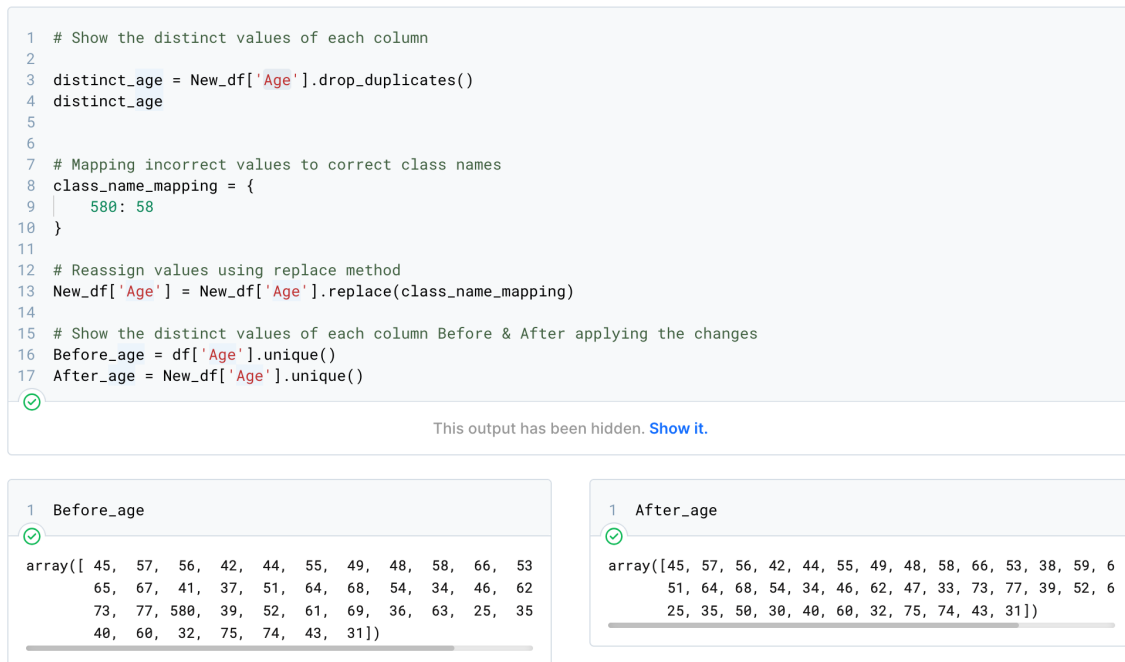


Figure 2.12: Error checking - typos

2.2.3 Zero versus null

- *Null values*³

One way for addressing null values involves examining various formats for representing null and then standardizing and replacing them with the term None.

The screenshot shows a Jupyter Notebook interface. The top cell contains Python code to check for null values in a specific column ('Nodes') and display the rows where missing values were found. The code is as follows:

```
1 # Check for null values in a specific column (e.g., 'column_name'), null values might be represented in a
2 # multiple formats (' ', null , naan )
3
4
5 # Check for null values in 'Nodes' column
6 missing_nodes_values = New_df['Nodes'].apply(lambda x: pd.isnull(x) or (isinstance(x, str) and (x.strip() == '' or x
7
8
9 # Display the rows where the column had missing values
10 New_df[missing_nodes_values]
```

The output of the code is a DataFrame with the following columns: Age (int64), Menopause (object), Tumor-size (object), Nodes (float64), Node-caps (object), Degree-of-malign... (object), Breast (object), and Recurrence (object). The first row shows a missing value in the 'Nodes' column (nan).

The second cell contains code to replace null values and empty strings with None:

```
1 #Replace null values and empty strings with None
2 New_df.loc[missing_nodes_values, 'Nodes'] = None
3
4 #Display the rows where the column had missing values
5 New_df[missing_nodes_values]
```

The output of this code is a warning message from pandas:

```
/shared-libs/python3.9/py/lib/python3.9/site-packages/pandas/core/indexing.py:1720: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-vers
self._setitem_single_column(loc, value, pi)
```

The warning message is followed by a DataFrame showing the result of the replacement. The 'Nodes' column now contains 'no' instead of 'nan'.

Figure 2.13: Handle null values

³The remaining features are added in the Appendices

2.2.4 Validation step




Figure 2.14: Validation step

Based on the report outcomes following the implementation of the necessary checks, it appears that additional duplicate rows have surfaced. Furthermore, we have identified an increased number of missing values in the dataset.

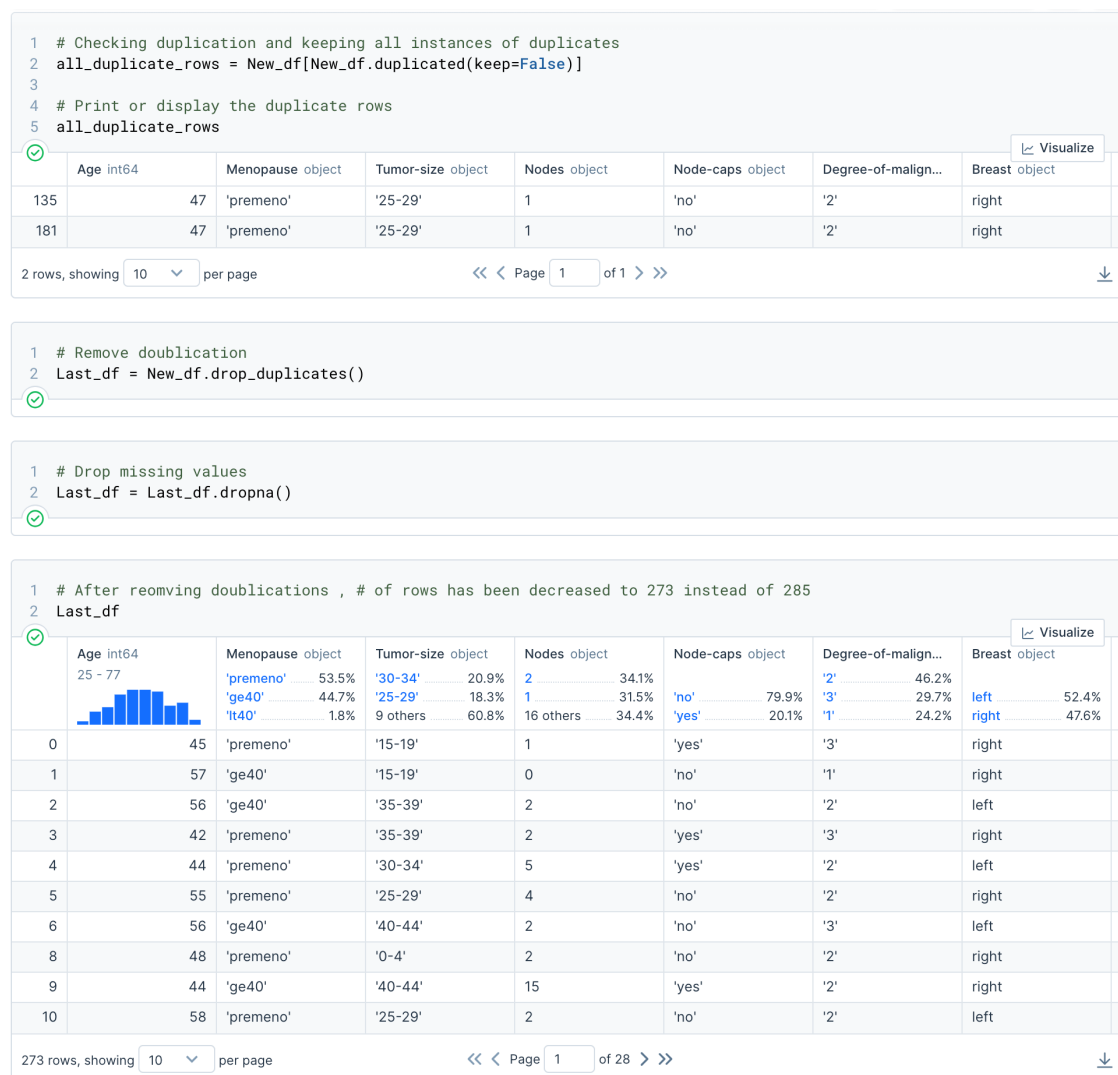


Figure 2.15: Resolving dataset after validation

Univariate Bivariate Analysis

3.1 Univariate Analysis

In data analysis, the initial step involves looking at each important variable individually. Univariate graphs show the distribution of data for a single variable, which could be a category like Menopause or a quantity like Age.

3.1.1 Qualitative variable

- *Menopause*¹

The graph indicates that the occurrence of symptoms is least frequent in premenopausal women, followed by menopausal women, and then postmenopausal women. This implies that the transition through menopause itself is linked to an increase in symptoms, and the frequency of symptoms remains heightened even after a woman has entered the postmenopausal stage.




Figure 3.1: Univariate Analysis - Menopause

¹The remaining univariate features are added in the Appendices

- *Breast*

The graph reveals that breast pain is more prevalent in the left breast compared to the right breast.




Figure 3.2: Univariate Analysis - Breast

3.1.2 Quantitative variable

- *Age*²

The graph indicates a right-skewed age distribution, implying a higher concentration of individuals in younger age groups compared to older ones. The median age stands at 50, signifying that half of the sample is below 50, and the other half is above 50.




Figure 3.3: Univariate Analysis - Age

²The remaining univariate features are added in the Appendices

- *Nodes*

This suggests that the highest number of nodes is concentrated between 0 and 5.

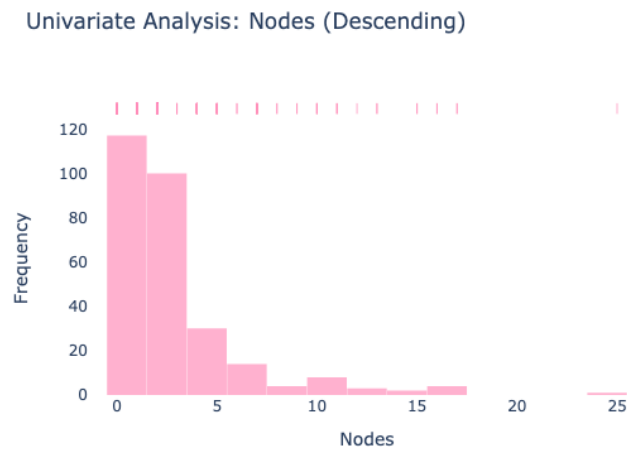


Figure 3.4: Univariate Analysis - Nodes

3.2 Bivariate Analysis

In research, a key question is understanding the relationship between two things, *A and B*. Bivariate graphs show this relationship by representing two variables together. The type of graph used depends on whether the variables are categories or quantities.

3.2.1 Categorical vs Quantitative

- *Menopause vs Age*³

The provided graph illustrates the distribution of women in a sample based on their menopausal status at different ages. It reveals an increasing proportion of menopausal women with age. Specifically, at the age of 40, approximately 20% of women are in the menopausal stage, and by the age of 50, this proportion rises to around 50%. As women reach the age of 60, nearly all have undergone menopause.




Figure 3.5: Bivariate Analysis - Menopause vs Age

³The remaining bivariate features are added in the Appendices

- *Menopause vs Tumor-size-encoded*

The bivariate analysis comparing menopause and tumor-size-encoded reveals a negative correlation between the two variables. This implies that menopausal women are less likely to have large tumors compared to women who are not in the menopausal stage.

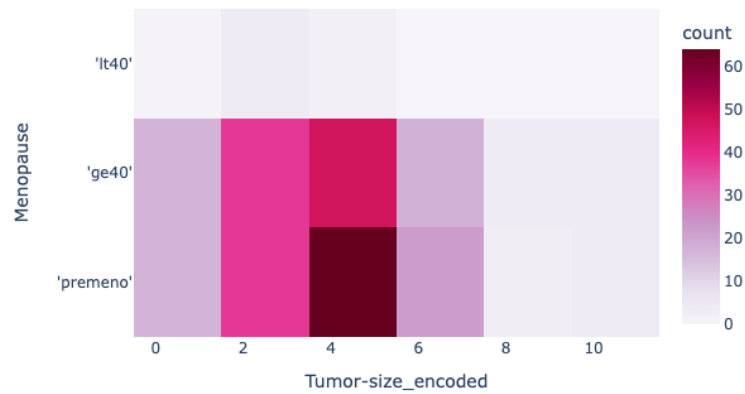


Figure 3.6: Bivariate Analysis - Menopause vs Tumor-size-encoded

3.2.2 Quantitative vs Quantitative

- *Age vs Nodes*

The bivariate analysis indicates that there might be an association between age and an increased risk of having more nodes.

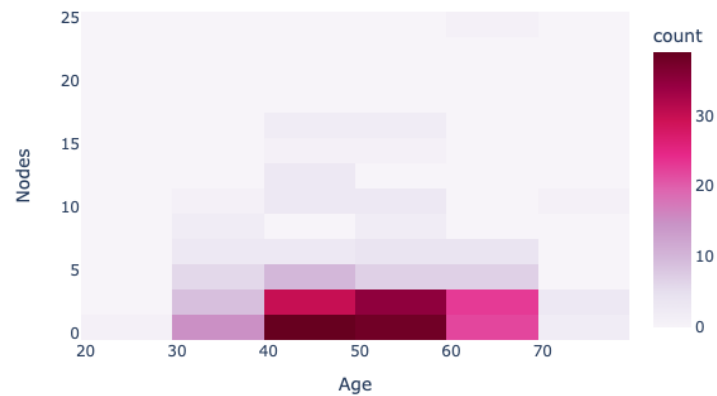


Figure 3.7: Bivariate Analysis - Age vs Node

- *Age vs Tumor-size_{encoded}*

The bivariate analysis reveals a positive correlation between the two variables, indicating that older individuals are more likely to have larger tumors compared to younger individuals.

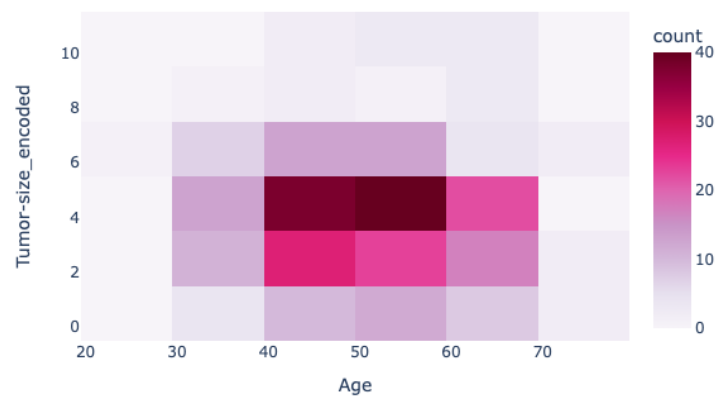


Figure 3.8: Bivariate Analysis - Age vs Tumor-size_{encoded}

Appendices

A Full Report

Data Management and Visualization Project

Sawsan Daban - Alaa AlSharekh



```
# Import Libraries
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
! pip install openpyxl
```

```
# To read the data
RawNormal = 'Raw Normal.xlsx'
df = pd.read_excel(RawNormal)
```

```
# Display the current data frame as df
df
```

	Age int64 25 - 580	Menopause object 'premeno' 52.2% 'ge40' 45% 2 others 2.8%	Tumor-size object '30-34' 21.1% '25-29' 19% 9 others 59.9%	nodes object 2 32.5% 1 31.5% 17 others 36%	Unnamed: 4 object 'no' 77.2% 'yes' 20.1% Missing 2.8%	degree-of-malign... '2' 45.7 '3' 29.8 '1' 24.6
0	45	'premeno'	'15-19'	1	'yes'	'3'
1	57	'ge40'	'15-19'	0	'no'	'1'
2	56	'ge40'	'35-39'	2	'no'	'2'
3	42	'premeno'	'35-39'	2	'yes'	'3'
4	44	'premeno'	'30-34'	5	'yes'	'2'
5	55	'premeno'	'25-29'	4	'no'	'2'
6	56	'ge40'	'40-44'	2	'no'	'3'
7	49	'premeno'	'10-14'		'no'	'2'
8	48	'premeno'	'0-4'	2	'no'	'2'
9	44	'ge40'	'40-44'	15	'yes'	'2'

289 rows, showing 10 per page << < Page 1 of 29 > >>

Data Preparation

Data summary

Population: The dataset appears to represent a sample of individuals, likely patients, from a medical context. These individuals are described based on various characteristics related to medical conditions.

Observations: Each row in the dataset represents an observation of an individual patient. The dataset has multiple observations, each corresponding to a different patient.

Dataset

Quantitative features: 1. Age: Numerical data representing the age of individuals. 2. Tumor-size: Although it appears as a string ('15-19', '35-39', etc.), it represents numerical intervals. 3. Nodes: Numerical data representing the number of nodes. It is a discrete numerical variable.

Categorical features: 1. Menopause: Categorical data representing the menopausal status of individuals ('premeno', 'ge40'). 2. Degree-of-malignance: Categorical data representing the degree of malignancy ('3', '1', '2'). 3. Breast: Categorical data indicating the side of the breast ('right' or 'left'). 4. Breast-quad: Categorical data representing the quadrant of the breast ('left_up', 'central', etc.). 5. Irradiation: Categorical data indicating whether irradiation was done ('yes' or 'no'). 6. Recurrence: Categorical data indicating the occurrence of events ('recurrence-events' or 'no-recurrence-events').

Quality report

```
# Initial # of Columns vs Rows
df.shape

(289, 10)

# General Info about the features of the dataset
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 289 entries, 0 to 288
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0    Age                 289 non-null   int64
1    Menopause           289 non-null   object
2    Tumor-size          289 non-null   object
3    nodes               289 non-null   object
4    Unnamed: 4          281 non-null   object
5    degree-of-malignance 289 non-null   object
6    Breast              289 non-null   object
7    Breast-quad         288 non-null   object
8    Irradiation         289 non-null   object
9    reccurrence         289 non-null   object
dtypes: int64(1), object(9)
memory usage: 22.7+ KB

# Data Profiling EDA
!pip install ydata-profiling

from ydata_profiling import ProfileReport

# Generate the data profiling report
report1 = ProfileReport(df, title='Raw Normal Profiling - report1 ')
report1.to_file("Raw Normal Profiling - report1.html")

!pip install ipywidgets==8.1.1

! pip install --upgrade pip

report1
```

Overview

Dataset statistics

Number of variables	10
Number of observations	289
Missing cells	9
Missing cells (%)	0.3%
Duplicate rows	6
Duplicate rows (%)	2.1%
Total size in memory	22.7 KiB
Average record size in memory	80.4 B

Variable types

Numeric	1
Categorical	8

The provided summary offers a concise overview of the existing dataset and highlights certain noteworthy aspects. According to the description, our dataset comprises 10 variables and encompasses a total of 289 observations. Among these variables, 8 fall under the categorical category, while 1 is designated as numerical variables & nodes with unknown variable type. Additionally, the summary points out the presence of duplicate rows and an unnamed column as part of the dataset's characteristics.

Please find below the full report: <https://msc-science-in-computing-2023.github.io/DMV-Reports/report1/>

Data quality control

Every set of data needs to undergo a form of quality control, which includes verifying for errors and ensuring uniform consistency, this might include : 1. Data consistency 2. Error checking 3. Zero versus null

Data Consistency

Data type

```
# Convert specific columns from object to string
string_columns = ['Menopause', 'Tumor-size', 'Unnamed: 4', 'nodes', 'Breast', 'Breast-quad', 'Irradiat:
df[string_columns] = df[string_columns].astype(str)

# Convert 'nodes' column to numeric, handling None values
int_columns = ['nodes']
df[int_columns]['nodes'] = pd.to_numeric(df[int_columns]['nodes'], errors='coerce')
```

```
# To describe the continious featurrs Age
df.describe()
```

	Age float64	
cou...	289	
me...	53.88927336	
std	32.74705546	
min	25	
25%	45	

50%	50
75%	59
max	580

8 rows, showing 10 per page << < Page 1 of 1 > >>

Duplications

Initially, the dataset has 4 duplicate rows. We handle them by removing duplicates.

Checking duplication
duplicate_rows = df[df.duplicated()
duplicate_rows

	Age int64	Menopause object	Tumor-size object	nodes object	Unnamed: 4 object	degree-of-malign...
64	46	'premeno'	'30-34'	16	'yes'	'3'
123	53	'premeno'	'25-29'	1	'yes'	'2'
157	60	'ge40'	'10-14'	1	'no'	'2'
242	56	'ge40'	'40-44'	7	'yes'	'3'

4 rows, showing 10 per page << < Page 1 of 1 > >>

Remove doublication
New_df = df.drop_duplicates()

After reomvng doublications , # of rows has been decreased to 285 instead of 289
New_df

	Age int64	Menopause object	Tumor-size object	nodes object	Unnamed: 4 object	degree-of-malign...
	25 - 580	'premeno' 52.3% 'ge40' 44.9% 2 others 2.8%	'30-34' 21.1% '25-29' 18.9% 9 others 60%	2 33% 1 31.2% 17 others 35.8%	'no' 77.9% 'yes' 19.3% nan 2.8%	'2' 45.6 '3' 29.5 '1' 24.9
0	45	'premeno'	'15-19'	1	'yes'	'3'
1	57	'ge40'	'15-19'	0	'no'	'1'
2	56	'ge40'	'35-39'	2	'no'	'2'
3	42	'premeno'	'35-39'	2	'yes'	'3'
4	44	'premeno'	'30-34'	5	'yes'	'2'
5	55	'premeno'	'25-29'	4	'no'	'2'
6	56	'ge40'	'40-44'	2	'no'	'3'
7	49	'premeno'	'10-14'		'no'	'2'
8	48	'premeno'	'0-4'	2	'no'	'2'
9	44	'ge40'	'40-44'	15	'yes'	'2'

285 rows, showing 10 per page << < Page 1 of 29 > >>

Dataset size has been decreased from 289 to 285.

New_df.shape
(285, 10)

Missing Column Name

As observed, it was noted that one column lacked a name. In response to understanding the problem statement and the values associated with this feature, a decision was made to rename it to Node-caps.

Assign a name to the unnamed column (Unnamed: 4)

New_df.rename(columns={'Unnamed: 4': 'Node-caps'}, inplace=True)

An additional step has been added, by unifying the column to start with capital letters.

Modify column names to start with Capital letters

New_df.rename(columns={'nodes': 'Nodes'}, inplace=True)

```
New_df.rename(columns={'degree-of-malignance': 'Degree-of-malignance'}, inplace=True)
New_df.rename(columns={'reccurrence': 'Reccurrence'}, inplace=True)
```

```
# After removing douplication & rename the columns
New_df
```

	Age int64 25 - 580	Menopause object 'premeno' 52.3% 'ge40' 44.9% 2 others 2.8%	Tumor-size object '30-34' 21.1% '25-29' 18.9% 9 others 60%	Nodes object 2 33% 1 31.2% 17 others 35.8%	Node-caps object 'no' 77.9% 'yes' 19.3% nan 2.8%	Degree-of-malign...
0	45	'premeno'	'15-19'	1	'yes'	'3'
1	57	'ge40'	'15-19'	0	'no'	'1'
2	56	'ge40'	'35-39'	2	'no'	'2'
3	42	'premeno'	'35-39'	2	'yes'	'3'
4	44	'premeno'	'30-34'	5	'yes'	'2'
5	55	'premeno'	'25-29'	4	'no'	'2'
6	56	'ge40'	'40-44'	2	'no'	'3'
7	49	'premeno'	'10-14'		'no'	'2'
8	48	'premeno'	'0-4'	2	'no'	'2'
9	44	'ge40'	'40-44'	15	'yes'	'2'

285 rows, showing 10 per page << < Page 1 of 29 > >>

Error Checking

Typos

To standardize the values, it was observed that certain entries contained typos.

```
# Show the distinct values of each column

distinct_breast = New_df['Breast'].drop_duplicates()
distinct_breast

# Mapping incorrect values to correct class names
class_name_mapping = {
    "right": "right",
    "left": "left",
    "rite": "right",
    "right'": "right",
    "rightt": "right",
    "'right'": "right",
    "'left'": "left",
    'right': "right"
}

# Reassign values using replace method
New_df['Breast'] = New_df['Breast'].replace(class_name_mapping)

# Show the distinct values of each column Before & After applying the changes
Before_Breast = df['Breast'].unique()
After_Breast = New_df['Breast'].unique()
```

Before_Breast

```
array(['right', 'left', 'rite', 'right', 'ri
```

After_Breast

```
array(['right', 'left'], dtype=object)
```

```
# Show the distinct values of each column

distinct_Reccurrence = New_df['Reccurrence'].drop_duplicates()
distinct_Reccurrence

# Mapping incorrect values to correct class names
class_name_mapping = {
    "recurrence-events": "recurrence-events",
    "no-recurrence-events": "no-recurrence-events",
    "'recurrence-events'" : "recurrence-events",
    "'no-recurrence-events'" : "no-recurrence-events",
    "R E": "recurrence-events"
}

}
```



```
# Reassign values using replace method
New_df['Reccurrence'] = New_df['Reccurrence'].replace(class_name_mapping)

# Show the distinct values of each column Before & After applying the changes
Before_Reccurrence = df['reccurrence'].unique()
After_Reccurrence = New_df['Reccurrence'].unique()
```

Before_Reccurrence

```
array(['recurrence-events', 'no-recurrence-even',
      dtype=object])
```

After_Reccurrence

```
array(['recurrence-events', 'no-recurrence-events',
      dtype=object])
```

```
# Show the distinct values of each column

distinct_Irradiation = New_df['Irradiation'].drop_duplicates()
distinct_Irradiation

# Mapping incorrect values to correct class names
class_name_mapping = {
    "no": "no",
    "N": "no",
    "yes": "yes",
    "'yes'" : "yes",
    "'no'" : "no"
}

# Reassign values using replace method
New_df['Irradiation'] = New_df['Irradiation'].replace(class_name_mapping)

# Show the distinct values of each column Before & After applying the changes
Before_Irradiation = df['Irradiation'].unique()
After_Irradiation = New_df['Irradiation'].unique()
```

Before_Irradiation

```
array(['no', 'yes', ' ', 'N'], dtype=object)
```

After_Irradiation

```
array(['no', 'yes', ' ', dtype=object])
```

Outlier

The average Age is approximately 52. Excluding the outlier with the value 580, the maximum Age is 77, and the minimum Age is 25. The data anomaly associated with the value 580 has been identified and corrected as shown below.

```
# Show the distinct values of each column
distinct_age = New_df['Age'].drop_duplicates()
distinct_age

# Mapping incorrect values to correct class names
class_name_mapping = {
    580: 58
}

# Reassign values using replace method
New_df['Age'] = New_df['Age'].replace(class_name_mapping)

# Show the distinct values of each column Before & After applying the changes
Before_age = df['Age'].unique()
After_age = New_df['Age'].unique()
```

Before_age

```
array([ 45,  57,  56,  42,  44,  55,  49,  48,  58,
        65,  67,  41,  37,  51,  64,  68,  54,  34,
        73,  77, 580,  39,  52,  61,  69,  36,  63,
        40,  60,  32,  75,  74,  43,  31])
```

After_age

```
array([45, 57, 56, 42, 44, 55, 49, 48, 58, 66, 53,
       51, 64, 68, 54, 34, 46, 62, 47, 33, 73, 77,
       25, 35, 50, 30, 40, 60, 32, 75, 74, 43, 31])
```

Validation step

```
from ydata_profiling import ProfileReport
```

```
# Generate the data profiling report to verify Data Consistency & Error Checking
report2 = ProfileReport(New_df, title='Raw Normal Profiling - report2')
report2.to_file("Raw Normal Profiling - report2.html")
```

report2

Overview

Dataset statistics

Number of variables	10
Number of observations	285
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	1
Duplicate rows (%)	0.4%
Total size in memory	24.5 KiB
Average record size in memory	88.0 B

Variable types

Numeric	1
Categorical	0

Please find below the full report: <https://msc-science-in-computing-2023.github.io/DMV-Reports/report2/>

Zero versus null

Null Values

Some techniques to handle Null values:

1. Removing Rows with Null Values. 2. Filling Null Values with a Specific Value like 'NA', 'Unknown' .. etc

```
# Check for null values in a specific column (e.g., 'column_name'), null values might be represented
# multiple formats (' ', null , naan )
```

```
# Check for null values in 'Menopause' column
missing_menopause_values = New_df['Menopause'].apply(lambda x: pd.isnull(x) or (isinstance(x, str) and x.isspace()))
```

```
# Display the rows where the column had missing values
New_df[missing_menopause_values]
```

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
36	46		'10-14'	0	'no'	'2'

1 row, showing 10 per page << < Page 1 of 1 > >>

```
# Replace null values and empty strings with None
New_df.loc[missing_menopause_values, 'Menopause'] = None

# Display the rows where the column had missing values
New_df[missing_menopause_values]
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self.setitem_single_column\(loc, value, pi\)](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self.setitem_single_column(loc, value, pi))

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
36	46	None	'10-14'	0	'no'	'2'

1 row, showing 10 per page << < Page 1 of 1 >>

```
array(['premeno', 'ge40', 'lt40', None], dtype=object)
```

```
# Display the rows where the column had missing values
New_df[missing_nodes_values]
```

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
7	49	'premeno'	'10-14'		'no'	'2'

1 row, showing 10 per page

[«](#)
[»](#)
Page 1 of 1

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
7	49	'premeno'	'10-14'	None	'no'	'2'

1 row, showing 10 per page
<< < Page 1 of 1 >>

```
array(['1', '0', '2', '5', '4', None, '15', '7', '3', '17', '10', '16',  
      '8', '6', '11', '9', '25', '13', '12'], dtype=object)
```

```
# Print the rows where the column has missing values
New_df[missing_nodescaps_values]
```

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
20	56	't40'	'20-24'	2	nan	'1'
31	68	'ge40'	'25-29'	3	nan	'1'

50	73	'ge40'	'15-19'	10	nan	'1'
54	48	'premeno'	'25-29'	1	nan	'2'
72	61	'ge40'	'25-29'	5	nan	'1'
93	51	'lt40'	'20-24'	0	nan	'1'
151	50	'ge40'	'30-34'	9	nan	'3'
267	57	'ge40'	'30-34'	11	nan	'3'

8 rows, showing 10 per page << < Page 1 of 1 > >>

Replace null values and empty strings with None
New_df.loc[missing_nodescaps_values, 'Node-caps'] = None

Display the rows where the column had missing values
New_df[missing_nodescaps_values]

/shared-lib/python3.9/py/lib/python3.9/site-packages/pandas/core/indexing.py:1720: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self_-setitem_single_column\(loc, value, pi\)](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self_-setitem_single_column(loc, value, pi))

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
20	56	'lt40'	'20-24'	2	None	'1'
31	68	'ge40'	'25-29'	3	None	'1'
50	73	'ge40'	'15-19'	10	None	'1'
54	48	'premeno'	'25-29'	1	None	'2'
72	61	'ge40'	'25-29'	5	None	'1'
93	51	'lt40'	'20-24'	0	None	'1'
151	50	'ge40'	'30-34'	9	None	'3'
267	57	'ge40'	'30-34'	11	None	'3'

< 8 rows, showing 10 per page << < Page 1 of 1 > >>

```
New_df['Node-caps'].unique()  
  
array(['yes', 'no', None], dtype=object)
```

Check for null values in a specific column (e.g., 'column_name'), null values might be represented
multiple formats (' ', null , naan)

Check for missing values in the 'Breast-quad' column
missing_breastquad_values = New_df['Breast-quad'].apply(lambda x: pd.isnull(x) or (isinstance(x, str) and x in [' ', 'naan']))

Print the rows where the column has missing values
New_df[missing_breastquad_values]

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
243	59	'ge40'	'30-34'	1	'no'	'3'

< 1 row, showing 10 per page << < Page 1 of 1 > >>

Replace null values and empty strings with None
New_df.loc[missing_breastquad_values, 'Breast-quad'] = None

Display the rows where the column had missing values
New_df[missing_breastquad_values]

/shared-lib/python3.9/py/lib/python3.9/site-packages/pandas/core/indexing.py:1720: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self_-setitem_single_column\(loc, value, pi\)](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self_-setitem_single_column(loc, value, pi))

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
243	59	'ge40'	'30-34'	1	'no'	'3'

```
New_df['Breast-quad'].unique()

array(['left_up', 'central', 'left_low', 'right_up',
       'right_low', None], dtype=object)
```

```
# Check for null values in a specific column (e.g., 'column_name'), null values might be represented
# multiple formats ( ' ', null , naan )

# Check for missing values in the 'Irradiation' column
missing_irradiation_values = New_df['Irradiation'].apply(lambda x: pd.isnull(x) or (isinstance(x, str) and x.strip() == ''))

# Print the rows where the column has missing values
New_df[missing_irradiation_values]
```

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
93	51	't40'	'20-24'	0	None	'1'

1 row, showing 10 per page

<< < Page 1 of 1 > >>

```
# Replace null values and empty strings with None
New_df.loc[missing_irradiation_values, 'Irradiation'] = None

# Display the rows where the column had missing values
New_df[missing_irradiation_values]
```

/shared-libs/python3.9/py/lib/python3.9/site-packages/pandas/core/indexing.py:1720: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self-setitem_single_column\(loc, value, pi\)](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-self-setitem_single_column(loc, value, pi))

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
93	51	't40'	'20-24'	0	None	'1'

1 row, showing 10 per page

<< < Page 1 of 1 > >>

```
New_df['Irradiation'].unique()

array(['no', 'yes', None], dtype=object)
```

New_df

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
	25 - 77	'premeno' 52.3% 2 others 47.4% Missing 0.4%	'30-34' 21.1% '25-29' 18.9% 9 others 60%	2 33% 17 others 66.7% Missing 0.4%	'no' 77.9% 'yes' 19.3% Missing 2.8%	'2' 45.6% '3' 29.5% '1' 24.9%
0	45	'premeno'	'15-19'	1	'yes'	'3'
1	57	'ge40'	'15-19'	0	'no'	'1'
2	56	'ge40'	'35-39'	2	'no'	'2'
3	42	'premeno'	'35-39'	2	'yes'	'3'
4	44	'premeno'	'30-34'	5	'yes'	'2'
5	55	'premeno'	'25-29'	4	'no'	'2'
6	56	'ge40'	'40-44'	2	'no'	'3'
7	49	'premeno'	'10-14'	None	'no'	'2'
8	48	'premeno'	'0-4'	2	'no'	'2'
9	44	'ge40'	'40-44'	15	'yes'	'2'

285 rows, showing 10 per page

<< < Page 1 of 29 > >>

Validation step

```
from ydata_profiling import ProfileReport

# Generate the data profiling report to verify Data Consistency & Error Checking
report3 = ProfileReport(New_df, title='Raw Normal Profiling - report3')
report3.to_file("Raw Normal Profiling - report3.html")
```

report3

Overview

Dataset statistics

Number of variables	10
Number of observations	285
Missing cells	12
Missing cells (%)	0.4%
Duplicate rows	1
Duplicate rows (%)	0.4%
Total size in memory	32.6 KiB
Average record size in memory	117.1 B

Variable types

Numeric	1
Categorical	0

Please find below the full report: <https://msc-science-in-computing-2023.github.io/DMV-Reports/report3/>

Based on the report outcomes following the implementation of the necessary checks, it appears that additional duplicate rows have surfaced. Furthermore, we have identified an increased number of missing values in the dataset.

```
# Checking duplication and keeping all instances of duplicates
all_duplicate_rows = New_df[New_df.duplicated(keep=False)]

# Print or display the duplicate rows
all_duplicate_rows
```

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
135	47	'premeno'	'25-29'	1	'no'	'2'
181	47	'premeno'	'25-29'	1	'no'	'2'

2 rows, showing 10 per page << < Page 1 of 1 > >>

```
# Remove dublication
Last_df = New_df.drop_duplicates()
```

```
# Drop missing values
Last_df = Last_df.dropna()
```

```
# After reomving dublications , # of rows has been decreased to 273 instead of 285
Last_df
```

	Age int64	Menopause object	Tumor-size object	Nodes object	Node-caps object	Degree-of-malign...
25 - 77		'premeno' 53.5%	'30-34' 20.9%	2 34.1%		'2' 46.2
		'ge40' 44.7%	'25-29' 18.3%	1 31.5%	'no' 79.9%	'3' 29.7

		't40' 1.8%	9 others 60.8%	16 others 34.4%	'yes' 20.1%	'1' 24.2%
0	45	'premeno'	'15-19'	1	'yes'	'3'
1	57	'ge40'	'15-19'	0	'no'	'1'
2	56	'ge40'	'35-39'	2	'no'	'2'
3	42	'premeno'	'35-39'	2	'yes'	'3'
4	44	'premeno'	'30-34'	5	'yes'	'2'
5	55	'premeno'	'25-29'	4	'no'	'2'
6	56	'ge40'	'40-44'	2	'no'	'3'
8	48	'premeno'	'0-4'	2	'no'	'2'
9	44	'ge40'	'40-44'	15	'yes'	'2'
10	58	'premeno'	'25-29'	2	'no'	'2'

273 rows, showing 10 per page << < Page 1 of 28 > >>

Last_df.shape

(273, 10)

Univariate & Bivariate analysis

Univariate Analysis

Qualitative variable

```
import plotly.graph_objects as go

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Menopause'

# Calculate the count of each category, including None values
column_counts = Last_df[column].value_counts(dropna=False).sort_index()

# Sort the values in descending order
column_counts_sorted = column_counts.sort_values(ascending=False)


# Define the custom bar color
custom_color = 'rgba(255, 144, 187, 1)'

# Create a bar chart
fig = go.Figure([go.Bar(x=column_counts_sorted.index, y=column_counts_sorted.values, marker_color=custom_color)])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column}',
                   xaxis_title=column,
                   yaxis_title='Frequency',
                   paper_bgcolor='white', # Set background color to white
                   plot_bgcolor='white', # Set plot background color to white
                   bargap=0.2, # Adjust the gap between bars
                   width=graph_width,
                   height=graph_height)

# Show the plot
fig.show()
```

Univariate Analysis: Menopause



```
import plotly.graph_objects as go

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Node-caps'

# Calculate the count of each category, including None values
column_counts = Last_df[column].value_counts(dropna=False).sort_index()

# Sort the values in descending order
column_counts_sorted = column_counts.sort_values(ascending=False)


# Define the custom bar color
custom_color = 'rgba(255, 144, 187, 1)'

# Create a bar chart
fig = go.Figure([go.Bar(x=column_counts_sorted.index, y=column_counts_sorted.values, marker_color=custom_color)])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column}',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.2, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()
```

Univariate Analysis: Node-caps




```

import plotly.graph_objects as go

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Breast'

# Calculate the count of each category, including None values
column_counts = Last_df[column].value_counts(dropna=False).sort_index()

# Sort the values in descending order
column_counts_sorted = column_counts.sort_values(ascending=False)

# Define the custom bar color
custom_color = 'rgba(255, 144, 187, 1)'


# Create a bar chart
fig = go.Figure([go.Bar(x=column_counts_sorted.index, y=column_counts_sorted.values, marker_color=custom_color)])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column}',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.2, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()

```

Univariate Analysis: Breast



```

import plotly.graph_objects as go

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Breast-quad'

# Calculate the count of each category, including None values
column_counts = Last_df[column].value_counts(dropna=False).sort_index()

# Sort the values in descending order
column_counts_sorted = column_counts.sort_values(ascending=False)

# Define the custom bar color
custom_color = 'rgba(255, 144, 187, 1)'

# Create a bar chart
fig = go.Figure([go.Bar(x=column_counts_sorted.index, y=column_counts_sorted.values, marker_color=custom_color)])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column}',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white

```


```

        plot_bgcolor='white', # Set plot background color to white
        bargap=0.2, # Adjust the gap between bars
        width=graph_width,
        height=graph_height)

# Show the plot
fig.show()

```

Univariate Analysis: Breast-quad



```

import plotly.graph_objects as go

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Irradiation'

# Calculate the count of each category, including None values
column_counts = Last_df[column].value_counts(dropna=False).sort_index()

# Sort the values in descending order
column_counts_sorted = column_counts.sort_values(ascending=False)

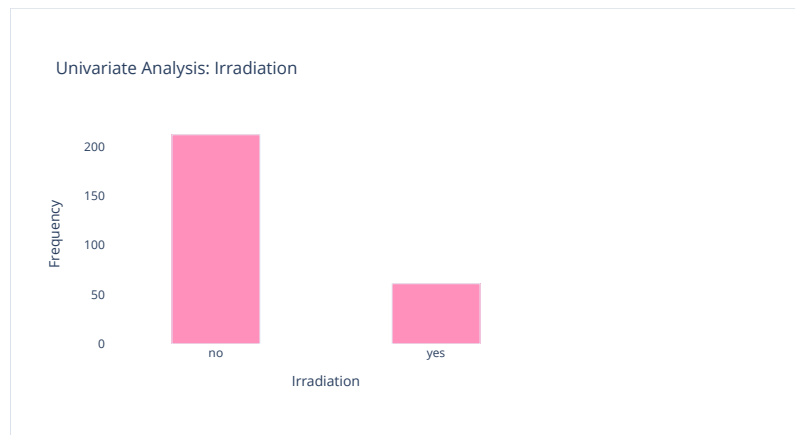
# Define the custom bar color
custom_color = 'rgba(255, 144, 187, 1)'

# Create a bar chart
fig = go.Figure([go.Bar(x=column_counts_sorted.index, y=column_counts_sorted.values, marker_color=custom_color)])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column}',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.2, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()

```



```
import plotly.graph_objects as go

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Reccurrence'

# Calculate the count of each category, including None values
column_counts = Last_df[column].value_counts(dropna=False).sort_index()


# Sort the values in descending order
column_counts_sorted = column_counts.sort_values(ascending=False)

# Define the custom bar color
custom_color = 'rgba(255, 144, 187, 1)'

# Create a bar chart
fig = go.Figure([go.Bar(x=column_counts_sorted.index, y=column_counts_sorted.values, marker_color=custom_color)])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column}',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.2, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()
```



Quantitative variable

```

import plotly.express as px

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Age'

# Sort the column in descending order
sorted_df = Last_df.sort_values(by=column, ascending=False)


# Create a histogram for the sorted column
fig = px.histogram(sorted_df, x=column, nbins=20, marginal='rug', color_discrete_sequence=['rgba(255, 0, 0, 0.5)'])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column} (Descending)',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.01, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()

```

Univariate Analysis: Age (Descending)



```

import plotly.express as px

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Tumor-size'

# Sort the column in descending order
sorted_df = Last_df.sort_values(by=column, ascending=False)


# Create a histogram for the sorted column
fig = px.histogram(sorted_df, x=column, nbins=20, marginal='rug', color_discrete_sequence=['rgba(255, 0, 0, 0.5)'])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column} (Descending)',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.01, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()

```

Univariate Analysis: Tumor-size (Descending)



```
import plotly.express as px

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
column = 'Nodes'


# Sort the column in descending order
sorted_df = Last_df.sort_values(by=column, ascending=False)

# Create a histogram for the sorted column
fig = px.histogram(sorted_df, x=column, nbins=20, marginal='rug', color_discrete_sequence=['rgba(255, 255, 255, 0.01)'])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column} (Descending)',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.01, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()
```

Univariate Analysis: Nodes (Descending)



```
import plotly.express as px

# Specify the size of the graph
graph_width = 600
graph_height = 400

# Select the column for analysis
```

```

column = 'Degree-of-malignance'

# Sort the column in descending order
sorted_df = Last_df.sort_values(by=column, ascending=False)


# Create a histogram for the sorted column
fig = px.histogram(sorted_df, x=column, nbins=20, marginal='rug', color_discrete_sequence=['rgba(255, 0, 0, 0.5)'])

# Update layout for better visualization and adjust the size
fig.update_layout(title_text=f'Univariate Analysis: {column} (Descending)',
                  xaxis_title=column,
                  yaxis_title='Frequency',
                  paper_bgcolor='white', # Set background color to white
                  plot_bgcolor='white', # Set plot background color to white
                  bargap=0.01, # Adjust the gap between bars
                  width=graph_width,
                  height=graph_height)

# Show the plot
fig.show()

```

Univariate Analysis: Degree-of-malignance (Descending)



Bivariate Analysis

We had to encode Tumor-size to a single definable numbers instead of an interval

```
from sklearn.preprocessing import LabelEncoder
```

```

# Create a copy to avoid SettingWithCopyWarning
Last_df = Last_df.copy()

# Encode 'Tumor-size' using label encoding
label_encoder = LabelEncoder()
Last_df['Tumor-size_encoded'] = label_encoder.fit_transform(Last_df['Tumor-size'])

```

Qualitative vs Quantitative


```

import plotly.express as px

# Using "PuRd" color scale with size 600x400
fig_purd = px.density_heatmap(
    Last_df, x="Age", y="Menopause",
    color_continuous_scale="PuRd",
    width=600, height=400
)

# Show the plots
fig_purd.show()


```



```
import plotly.express as px
from sklearn.preprocessing import LabelEncoder

# Using "PuRd" color scale with size 600x400
fig_purd = px.density_heatmap(
    Last_df, x="Tumor-size_encoded", y="Menopause",
    color_continuous_scale="PuRd",
    width=600, height=400
)


# Show the plots
fig_purd.show()
```



```
import plotly.express as px

# Using "PuRd" color scale with size 600x400
fig_purd = px.density_heatmap(
    Last_df, x="Menopause", y="Nodes",
    color_continuous_scale="PuRd",
    width=600, height=400
)

# Show the plots
fig_purd.show()
```



Qualitative vs Quantitative


```
import plotly.express as px

# Assuming Last_df is your DataFrame
# Replace "Last_df" with the actual name of your DataFrame

# Sort the values of the 'Tumor-size_encoded' column in descending order
Last_df_sorted = Last_df.sort_values(by='Tumor-size_encoded', ascending=False)

# Using "PuRd" color scale with size 600x400
fig_purd = px.density_heatmap(
    Last_df_sorted, x="Age", y="Tumor-size_encoded",
    color_continuous_scale="PuRd",
    width=600, height=400
)

# Show the plot
fig_purd.show()
```




```
import plotly.express as px

# Sort the values of the 'Tumor-size' column in descending order
Last_df_sorted = Last_df.sort_values(by=['Nodes'], ascending=False, na_position='first')

# Using "PuRd" color scale with size 600x400
fig_purd = px.density_heatmap(
    Last_df_sorted, x="Age", y="Nodes",
    color_continuous_scale="PuRd",
    width=600, height=400
)
```




```
# Show the plots
fig_purd.show()
```



```
import plotly.express as px

# Using "PuRd" color scale with size 600x400
fig_purd = px.density_heatmap(
    Last_df, x="Tumor-size_encoded", y="Nodes",
    color_continuous_scale="PuRd",
    width=600, height=400
)

# Show the plots
fig_purd.show()
```



Validation step

```
from ydata_profiling import ProfileReport

# Generate the data profiling report to verify Data Consistency & Error Checking
report4 = ProfileReport(Last_df, title='Raw Normal Profiling - report4')
report4.to_file("Raw Normal Profiling - report4.html")
```

report4

Overview

Dataset statistics	
Number of variables	11
Number of observations	273
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	25.6 KiB
Average record size in memory	96.0 B
Variable types	
Numeric	2
Categorical	9

Please find below the full report: <https://msc-science-in-computing-2023.github.io/DMV-Reports/report4/>